

# Assignment4

October 25, 2024

## 0.0.1 CS2101 - Programming for Science and Finance

Prof. Götz Pfeiffer School of Mathematical and Statistical Sciences University of Galway

---

## 1 Computer Lab 4

---

### 1.1 0. Image Manipulation

```
[1]: import numpy as np
      from PIL import Image
```

```
[2]: # Helper function to display image
      def display_image(im: np.ndarray, scale : int = 1) -> Image:
          imageHeight, imageWidth = im.shape[0], im.shape[1]
          return display(Image.fromarray(im).resize((imageHeight*scale,
↪imageWidth*scale)))
```

```
[3]: howl = np.asarray(Image.open('howl.ppm'))
      display_image(howl)
```



### 1.2 1. Inverting an image

- The first task is to write a function `invert` that takes a value  $i$  of type `uint8` and returns  $255 - i$ .
- If we apply this to a colour intensity it will invert the intensity: zero intensity becomes full intensity and vice versa.

```
[4]: def invert(i : np.uint8) -> np.uint8:
      # Invert the pixel value
      return np.uint8(255 - i)
```

- Test on a value of type np.uint8:

```
[5]: invert(np.uint8(100))
```

```
[5]: np.uint8(155)
```

- Next you can use the function `invert` to invert all the colour channels of an image.

```
[6]: def invert_image(im: np.ndarray) -> np.ndarray:
      # Create a copy of the image to avoid modifying the original in-place
      inverted_image = im.copy()

      # Iterate over each pixel and channel, apply the invert function
      for i in range(im.shape[0]):
          for j in range(im.shape[1]):
              inverted_image[i, j, 0] = invert(im[i, j, 0]) # Red channel
              inverted_image[i, j, 1] = invert(im[i, j, 1]) # Green channel
              inverted_image[i, j, 2] = invert(im[i, j, 2]) # Blue channel

      return inverted_image
```

```
[7]: display_image(invert_image(howl))
```



## 1.3 2. Converting an Image to Greyscale

- In this task you will write a function `greyscale` that takes an image and returns a greyscale version of it.
- This is done by letting every pixel have an equal intensity of red, green, and blue in the greyscale.
- That common intensity should be the **average** of the red, green, and blue in the original image.

```
[8]: def average(rgb):
      return np.mean(rgb)
```

```
[9]: def greyscale(im):
    greyScale = im.copy()

    for i in range(im.shape[0]):
        for j in range(im.shape[1]):
            greyScale[i][j] = average(im[i, j])

    return greyScale
```

```
[10]: display_image(greyscale(howl))
```



### 1.4 3. Quantizing an image

- If we use `uint8` as the type of the colour intensities, that means that there are  $256^3 = 2^{24} = 16777216$  different colours that a pixel can have.
- In this task, you'll write a function that reduces the quality of an image by grouping similar intensities together and giving them a common value.

#### 1.4.1 Quantizing a single value

- Your first subtask is to write a function `quantize` that takes a `uint8` value and a parameter `bits` that determines how many different 'bins' of values we create.
- For example: if `bits` has the value 6, there should be  $2^6 = 64$  bins, and we let the values 0,1,2,3 be grouped together to give the result 0, the values 4,5,6,7 should all return 4, and so on.
- If `bits` instead have the value 4, there should be  $2^4 = 16$  bins, and the values 0,1,...,15 should return 0; the values 16,17,...,31 should all return 16 etc.

```
[11]: def quantize(val : np.uint8, bits : int):# -> Any / int:
    if bits >= 8 or bits <= 0:
        return val
    else:
        numBins = 2 ** bits
        binSize = 256 / numBins
        return np.uint8((val // binSize) * binSize)
```

- Test:

```
[12]: print(quantize(17, 4))
print(quantize(17, 6))
```

16  
16

### 1.4.2 Quantizing an image

The next task is to write a function `quantize_image` that takes two arguments, the image to quantize, and again a `bits` parameter, that determines how coarsely we quantize the colours in the image.

```
[13]: def quantize_image(im, bits):  
    quantizedImage = im.copy()  
    # Apply the quantization on each color channel independently using  
    ↪vectorized operations  
    quantizedImage[:, :, 0] = quantize(im[:, :, 0], bits)  
    quantizedImage[:, :, 1] = quantize(im[:, :, 1], bits)  
    quantizedImage[:, :, 2] = quantize(im[:, :, 2], bits)  
  
    return quantizedImage
```

```
[14]: display_image(quantize_image(howl, 4))
```



## 1.5 4. Downsampling/Upsampling

### 1.5.1 Upsampling

By upsampling in here we mean that we take an original image, and where each pixel is replaced by a  $n \times n$  block of pixels in the new image. Each pixel in the block is identical to the original pixel.

```
[15]: def upsample(im, n):  
    upsampledImage = np.zeros((im.shape[0] * n, im.shape[1] * n, im.shape[2]),  
    ↪dtype=np.uint8)  
  
    for i in range(im.shape[0]):  
        for j in range(im.shape[1]):  
            upsampledImage[i * n:i * n + n, j * n:j * n + n] = im[i, j]  
  
    return upsampledImage
```

```
[16]: display_image(upsample(howl, 2))
```



### 1.5.2 Downsampling

- When we downsample an image `im`, we divide the image into blocks of size  $n \times n$ , and compute the average colour intensities in these blocks.
- Each such block then gives rise to a single pixel in the resulting image.
- Note that first upsampling with a factor `n` and then downsampling with the same factor should return an image identical to the original.
- Whereas first downsampling and then upsampling does not return an identical image.

```
[17]: def downsample(im, n):
        downsampledImage = np.zeros((im.shape[0] // n, im.shape[1] // n, im.
        ↪shape[2]), dtype=np.uint8)

        for i in range(im.shape[0] // n):
            for j in range(im.shape[1] // n):
                block = im[i * n:(i + 1) * n, j * n:(j + 1) * n]
                meanBlockColour = np.mean(block, axis=(0, 1))

                downsampledImage[i, j] = meanBlockColour
        return downsampledImage
```

```
[18]: display_image(downsample(howl, 2))
```



## 1.6 5. Noise

- In this last part of the assignment, you will write a function `noisify` to add noise to an image.
- By this we mean that we add a random number to each colour channel of each pixel.

- The call `noisify(im, sigma)` should add a normally distributed random number with mean value 0 and standard deviation `sigma` to each entry in the array representing the image `im`.
- Make sure that:
  1. The result after adding is converted to a `uint8`.
  2. The result is in the interval 0 to 255.

```
[19]: import random
def noisify(im, sigma):
    noiseArray = np.random.normal(0, sigma, im.shape)
    noisyImage = np.clip(im.astype(np.float32) + noiseArray, 0, 255)
    return noisyImage.astype(np.uint8)
```

```
[20]: display_image(noisify(howl, 4), scale=2)
display_image(noisify(howl, 16), scale=2)
```

